


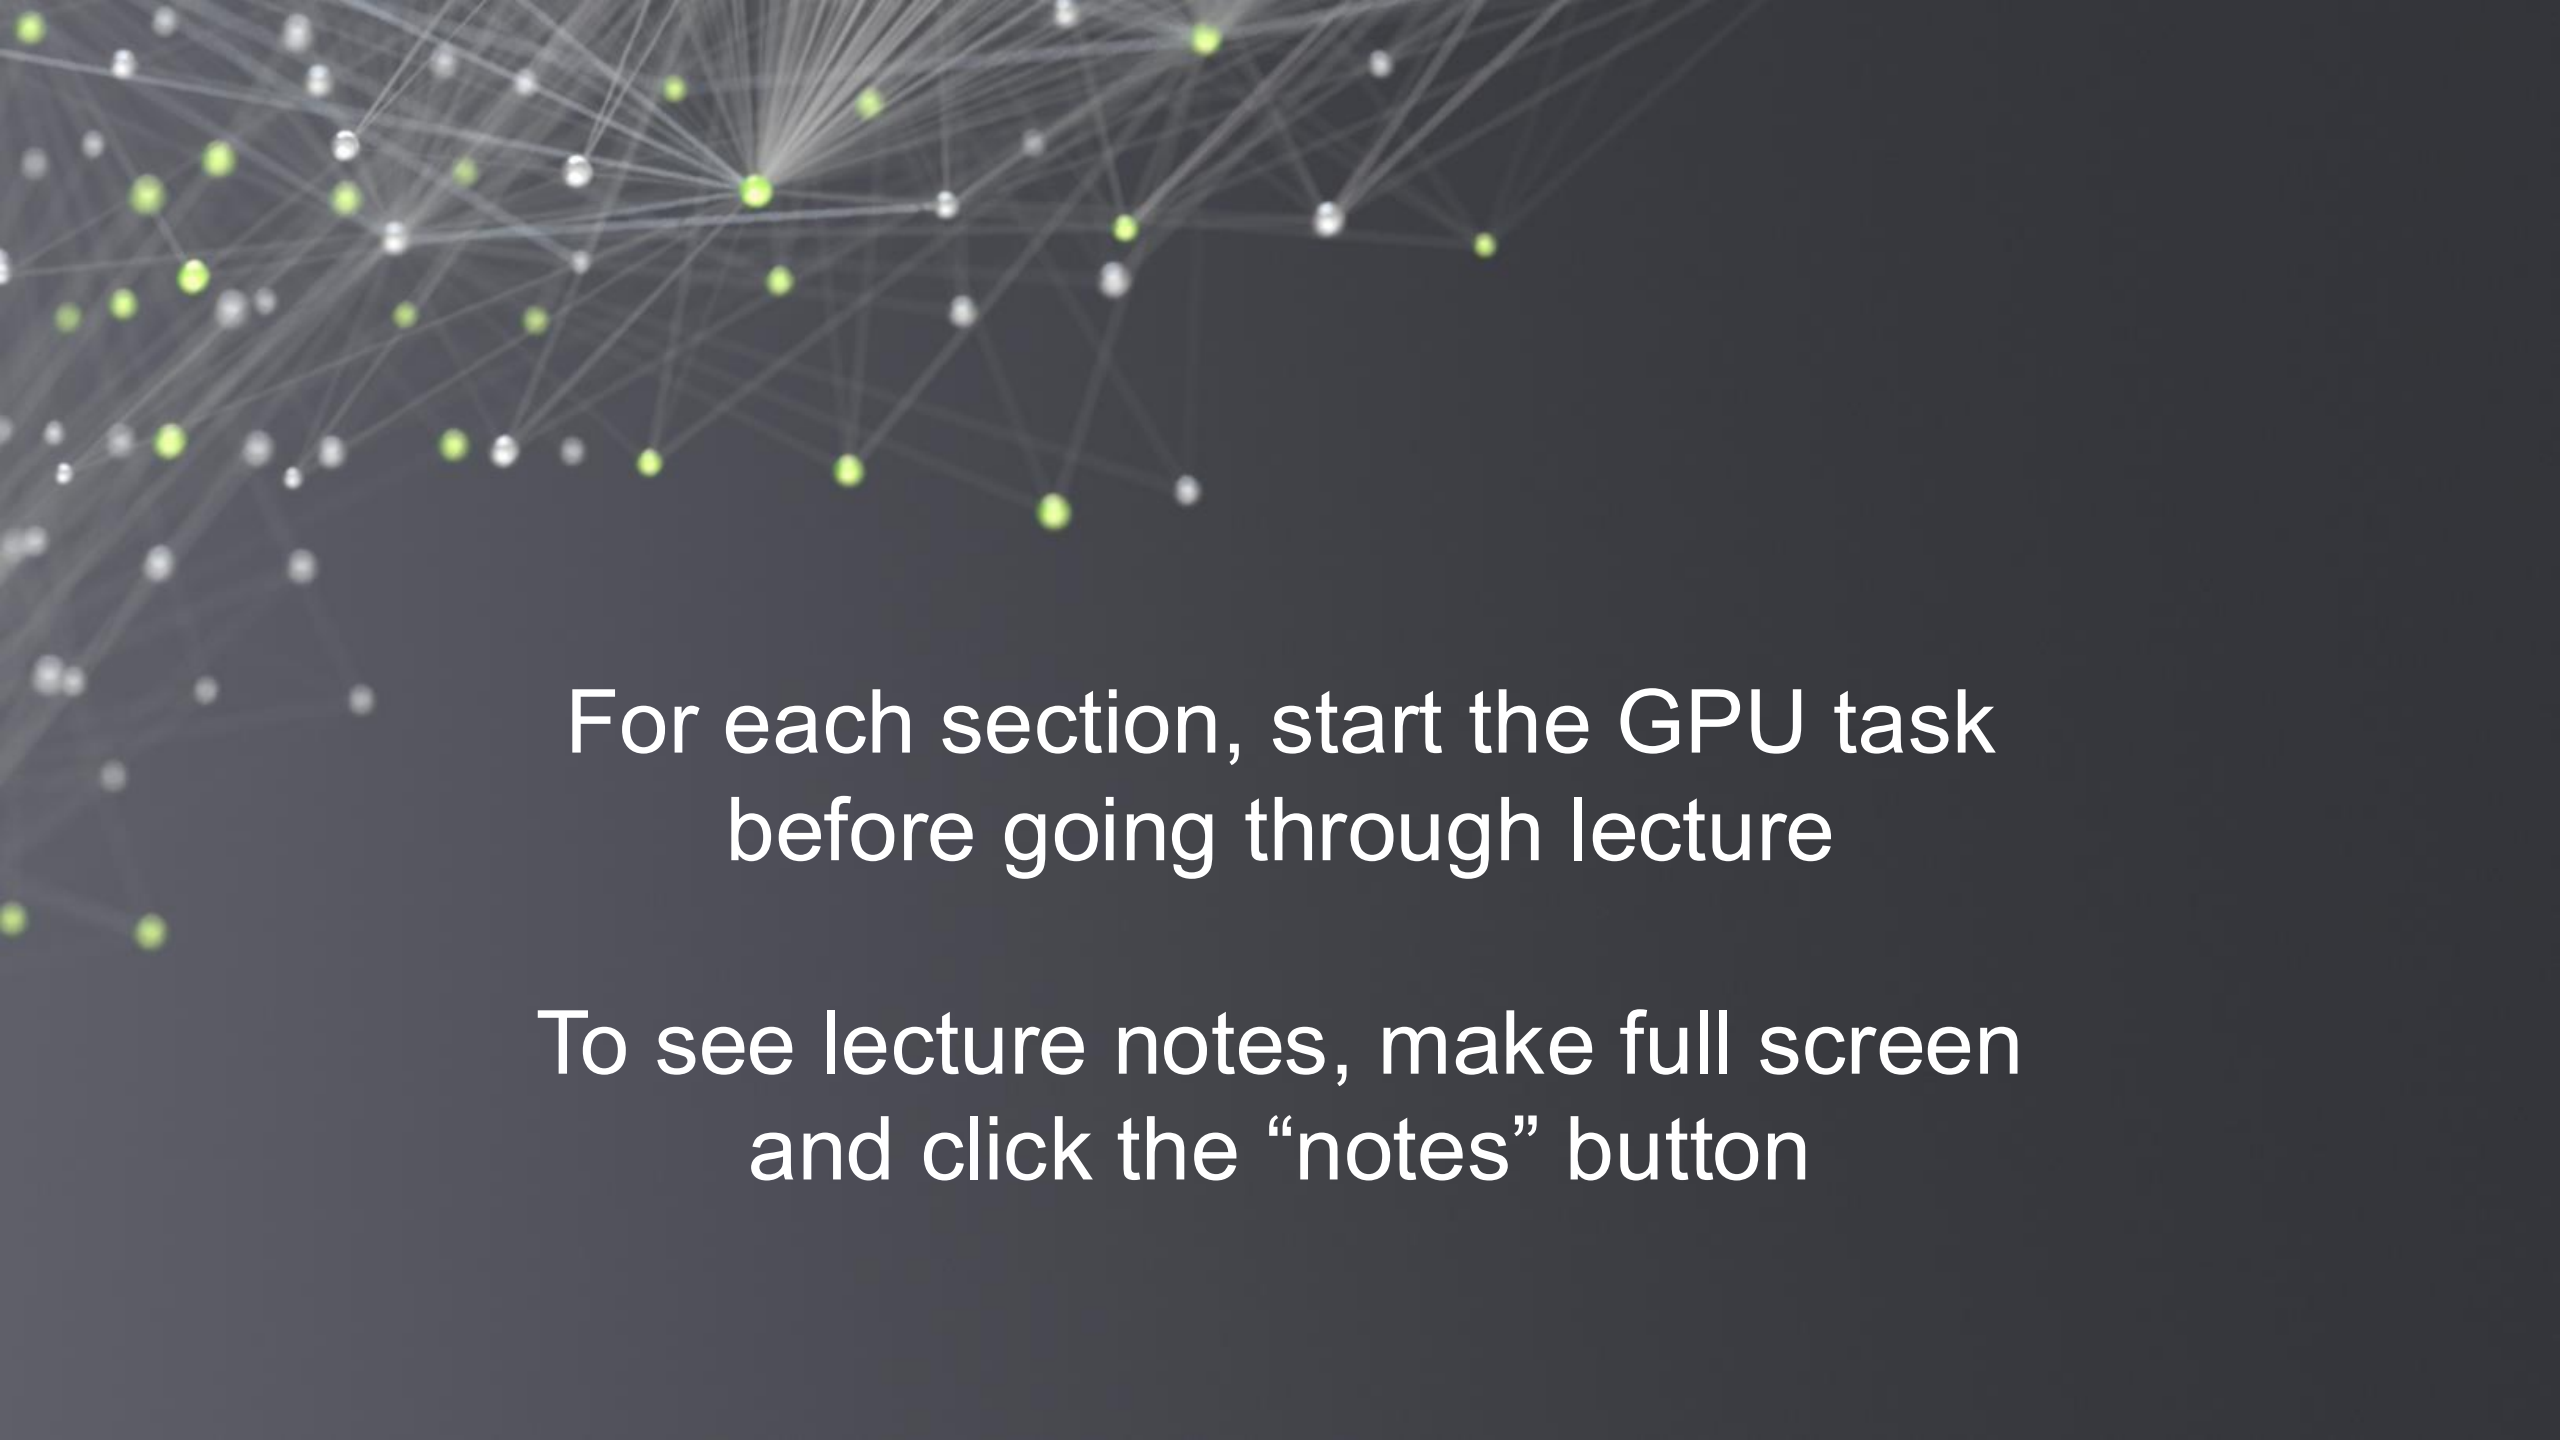


DEEP  
LEARNING  
INSTITUTE

# ACCELERATING DATA ENGINEERING PIPELINES

Part 1: Data Storage





For each section, start the GPU task  
before going through lecture

To see lecture notes, make full screen  
and click the “notes” button



**WELCOME!**



Main Goal:  
How do we organize and process an  
unexplored dataset to produce  
actionable insight?

# THE GOALS OF THIS COURSE

- Get used to many different data types / frameworks and how they operate on GPU vs CPU.
- Understand how DAG based frameworks can speed up ETL
- Learn how to visualize data to
  - Assess data quality
  - Allow users to make their own decisions through interactivity



---

# AGENDA

Part 1: Data Formats

Part 2: ETL with NVTabular

Part 3: Data Visualization

# AGENDA – PART I

- Systems Engineering
- File Formats
- Data Frameworks
- Lab




SYSTEMS ENGINEERING



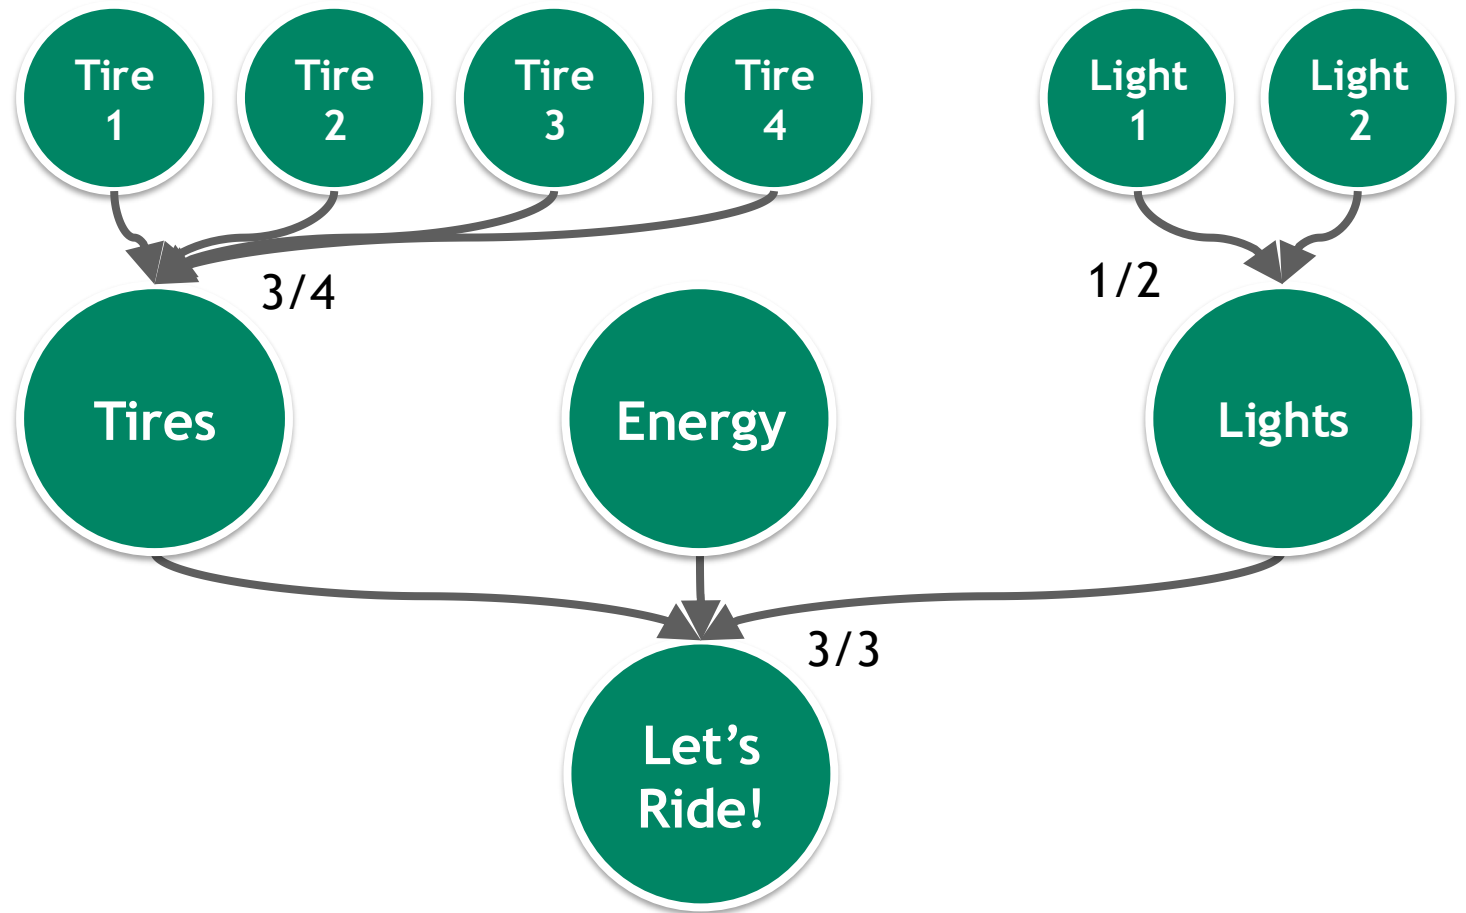
# SYSTEM EXAMPLE

## Modeling a Car



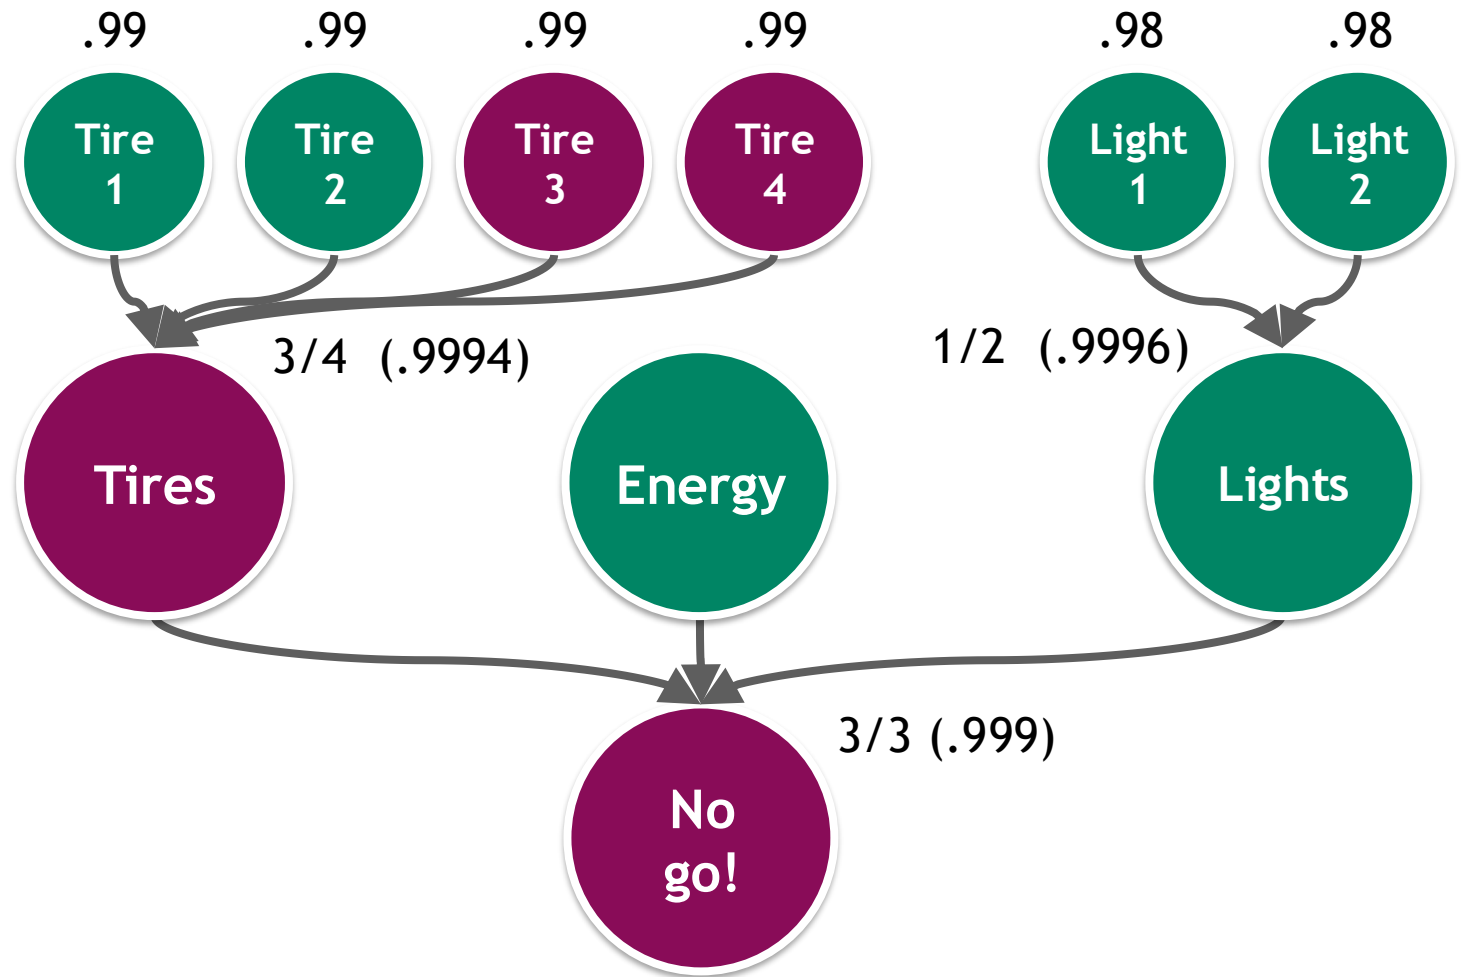
# SYSTEM EXAMPLE

## Modeling a Car



# SYSTEM EXAMPLE

## Modeling a Car






# SYSTEMS BIG AND SMALL

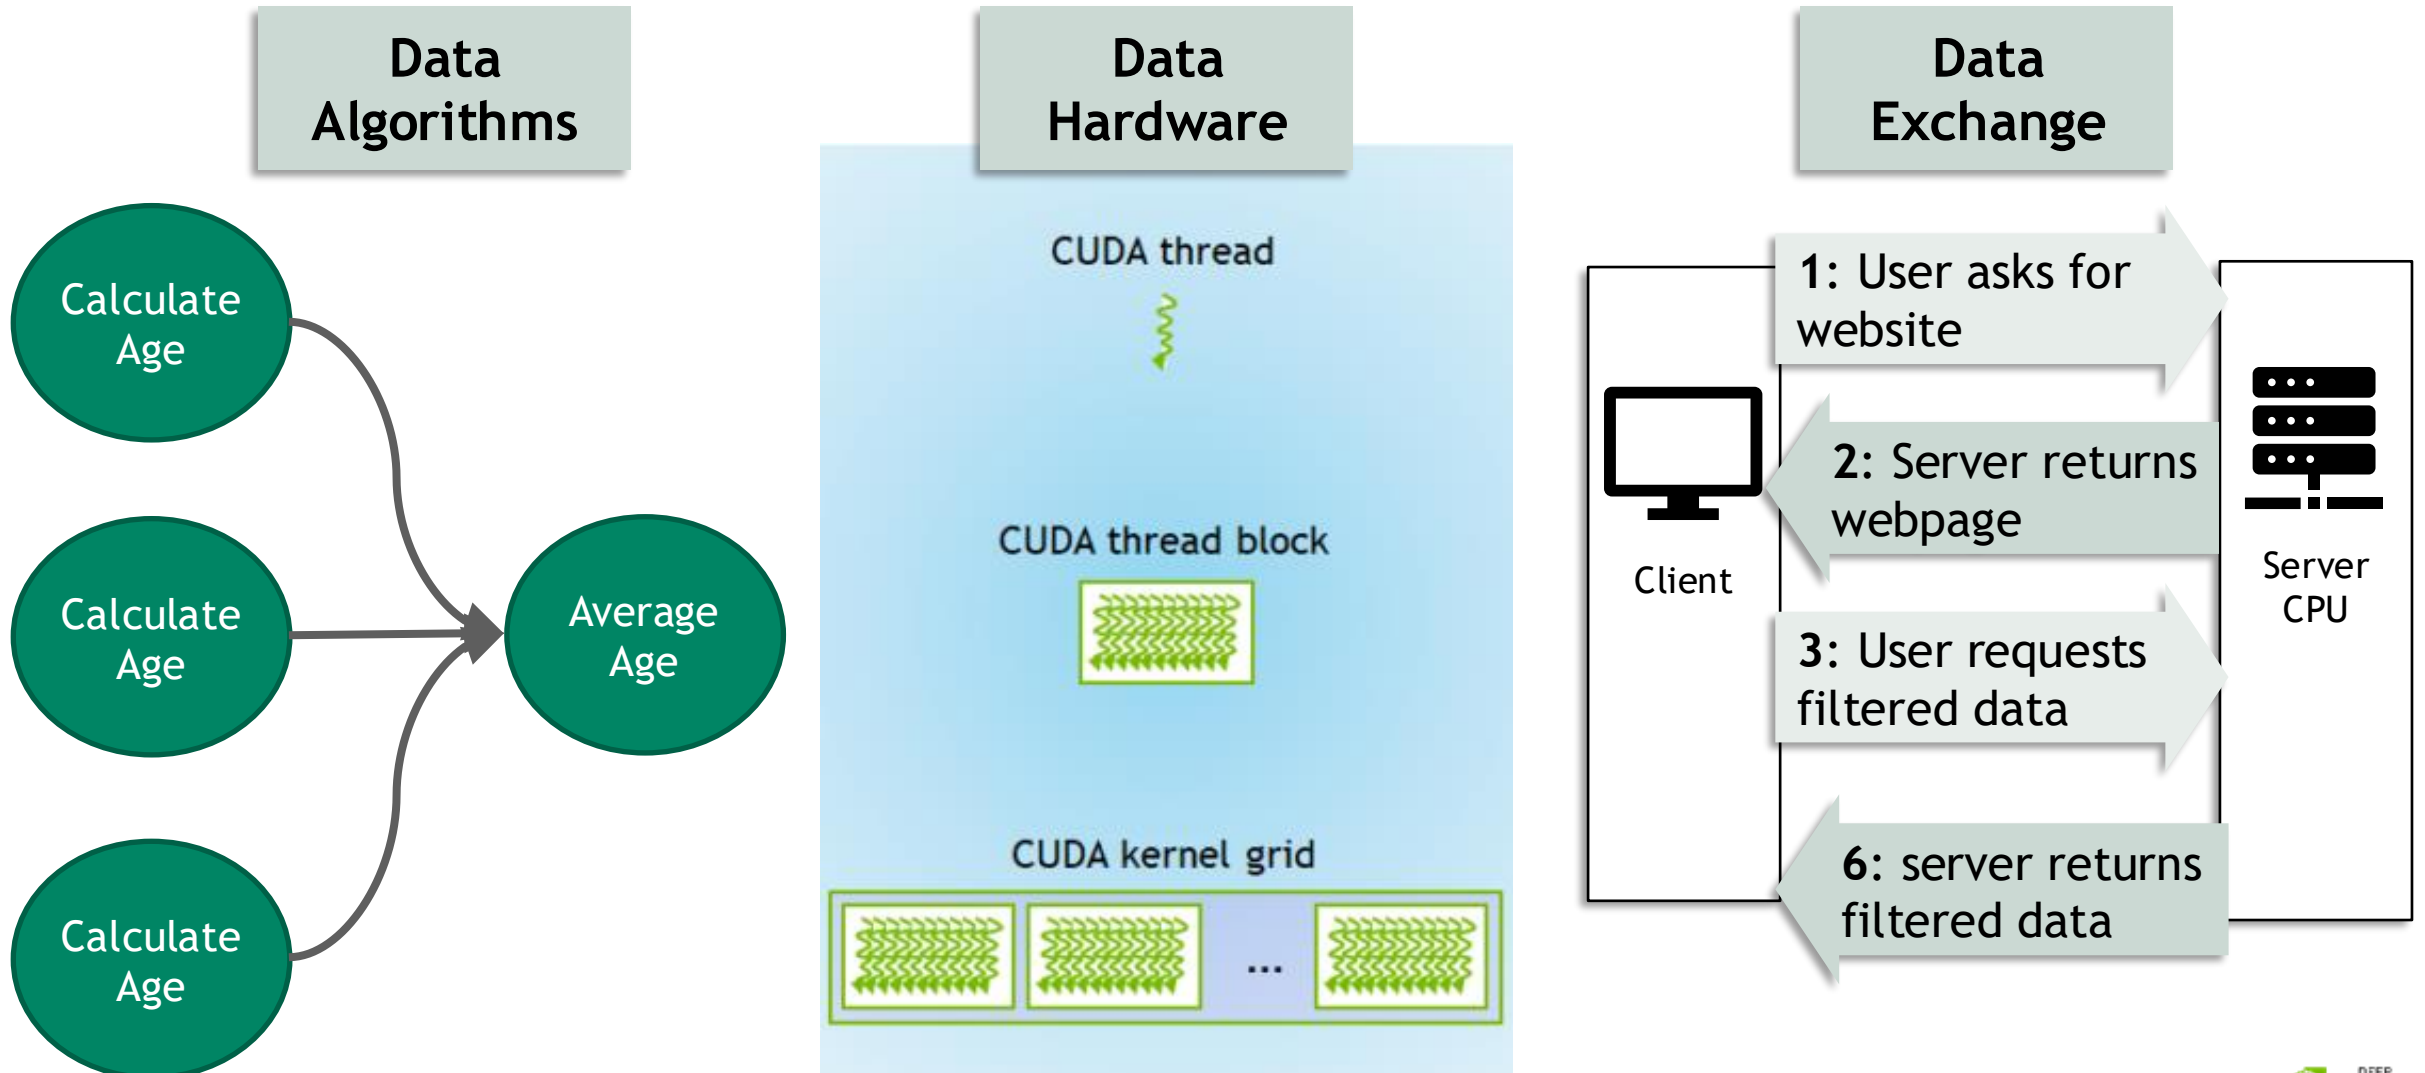
Details



Big Picture



# SYSTEM ENGINEERING FOR DATA



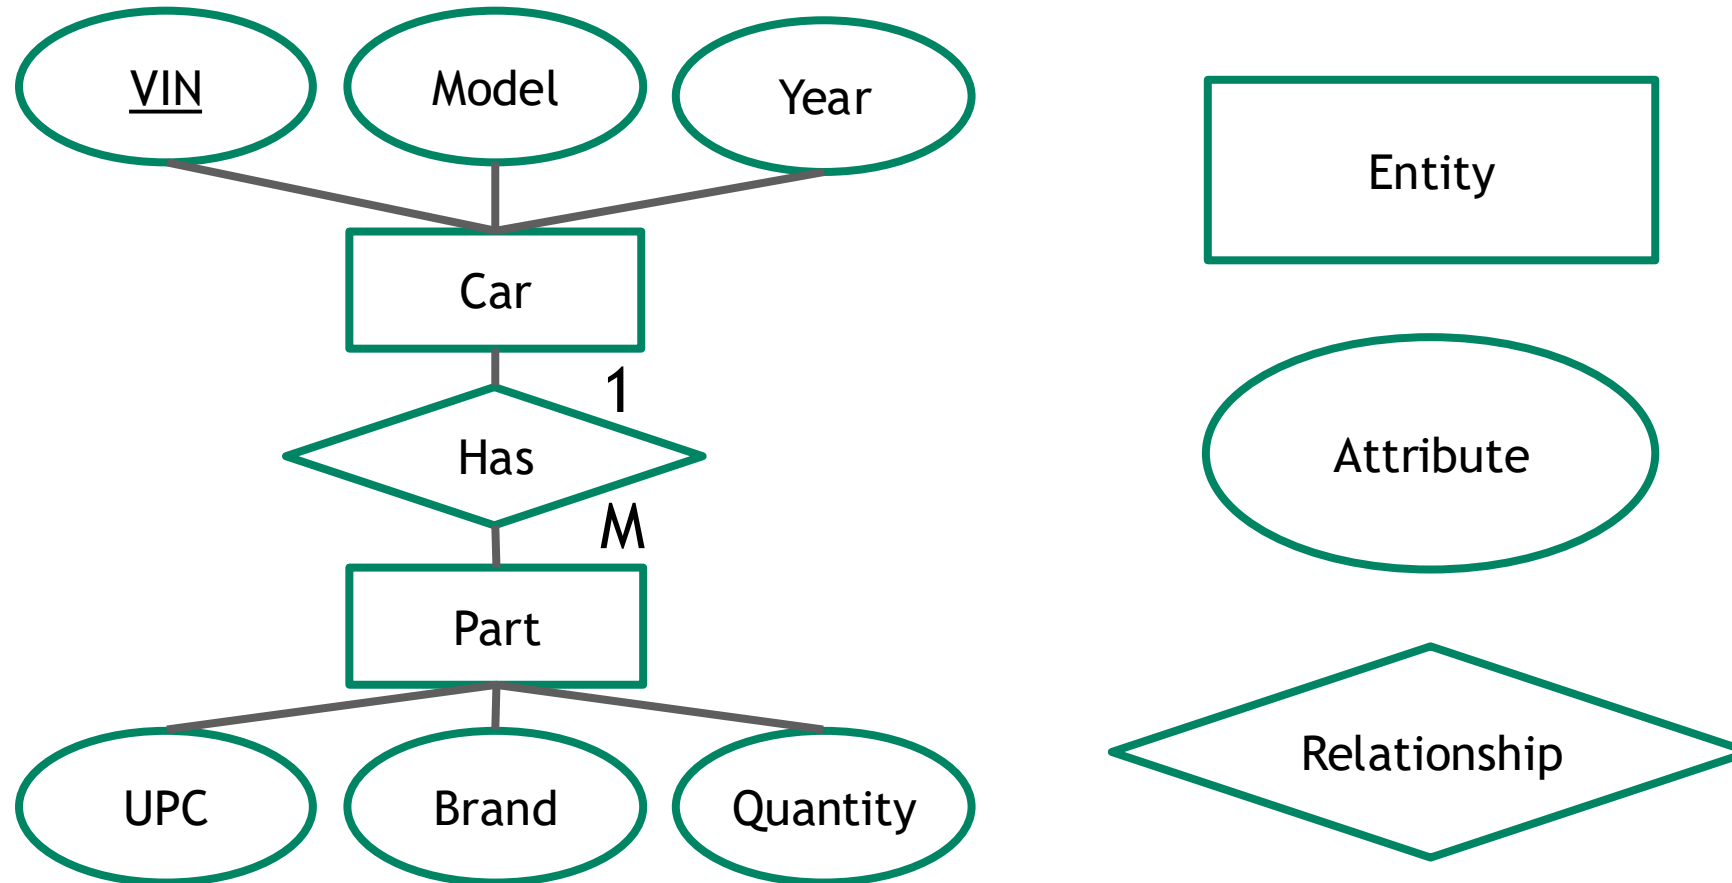




# DATA AS A SYSTEM OF ALGORITHMS

# DATABASE SYSTEM DESIGN

## Enhanced Entity Relationship Diagram



# DATABASE SYSTEM DESIGN

## From Design to Practice

### Cars.csv

VIN, Model, Year  
1a2b3c, Sedan, 1986  
4d5e6g, Convertible, 2011  
7h8i9j, Sedan, 1997

### Parts.csv

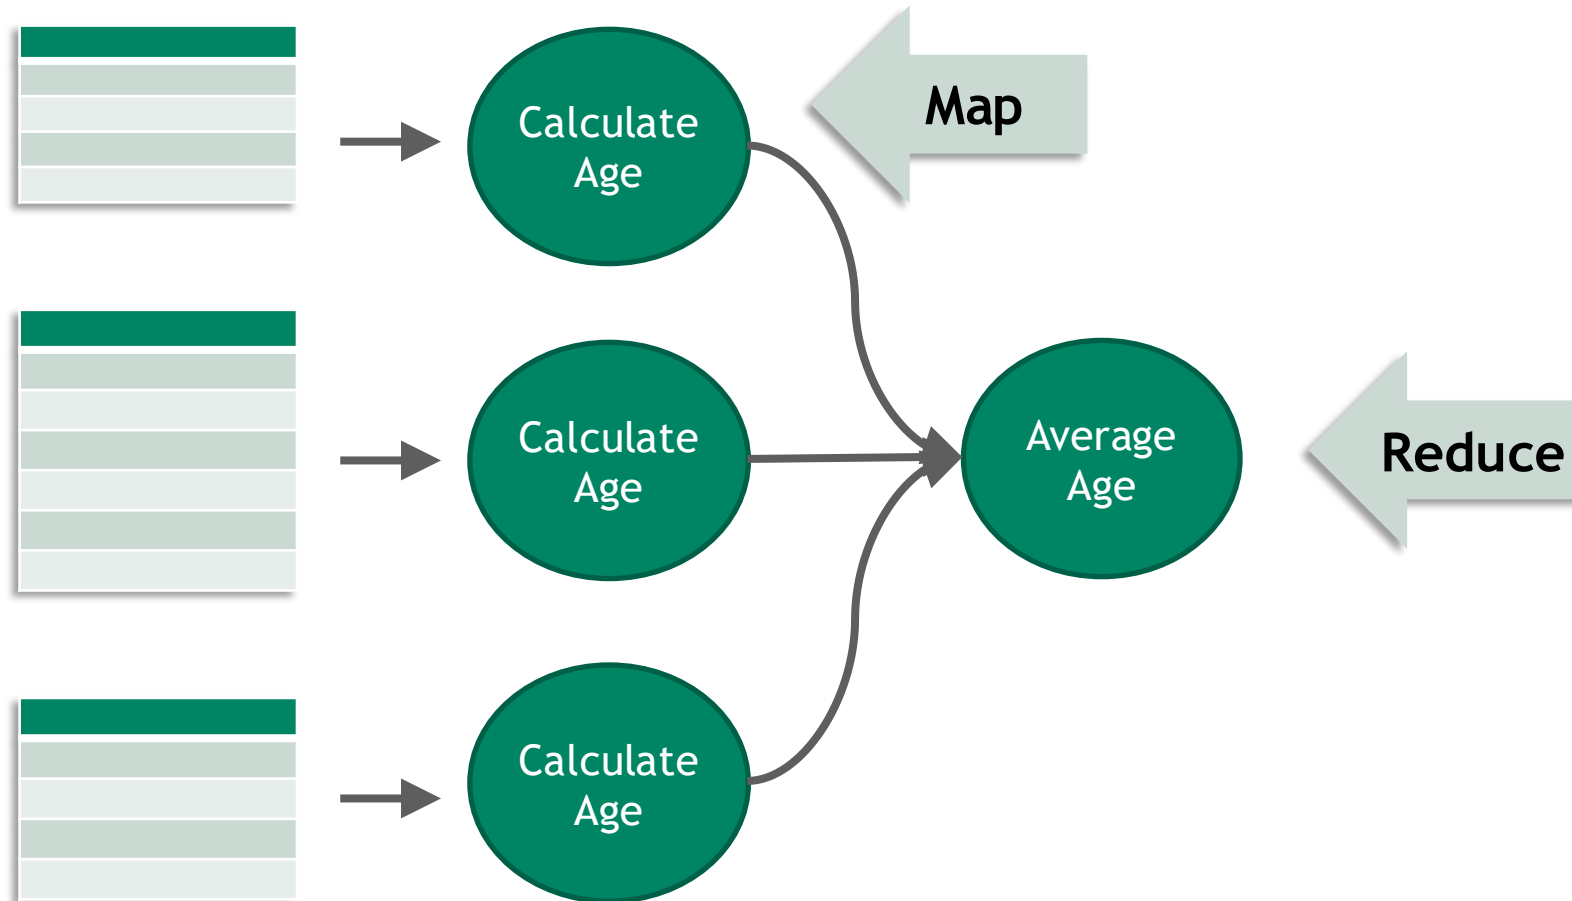
VIN, UPC, Brand, Quantity  
1a2b3c, 8675309, Generic Lights, 2  
1a2b3c, 8675310, Generic Tires, 4  
4d5e6g, 8675309, Awesome Lights, 2  
4d5e6g, 8675310, Awesome Tires, 4

### Cars.json

```
{  
  {  
    "VIN": 1a3b3c,  
    "Model": "Sedan",  
    "Year": 1986,  
    "Parts": [  
      {  
        "UPC": 8675309,  
        "Brand": "Generic Lights",  
        "Quantity": 2  
      }, {  
        "UPC": 8675310,  
        "Brand": "Generic Lights",  
        "Quantity": 4  
      }  
    ]  
  },  
  ...  
}
```

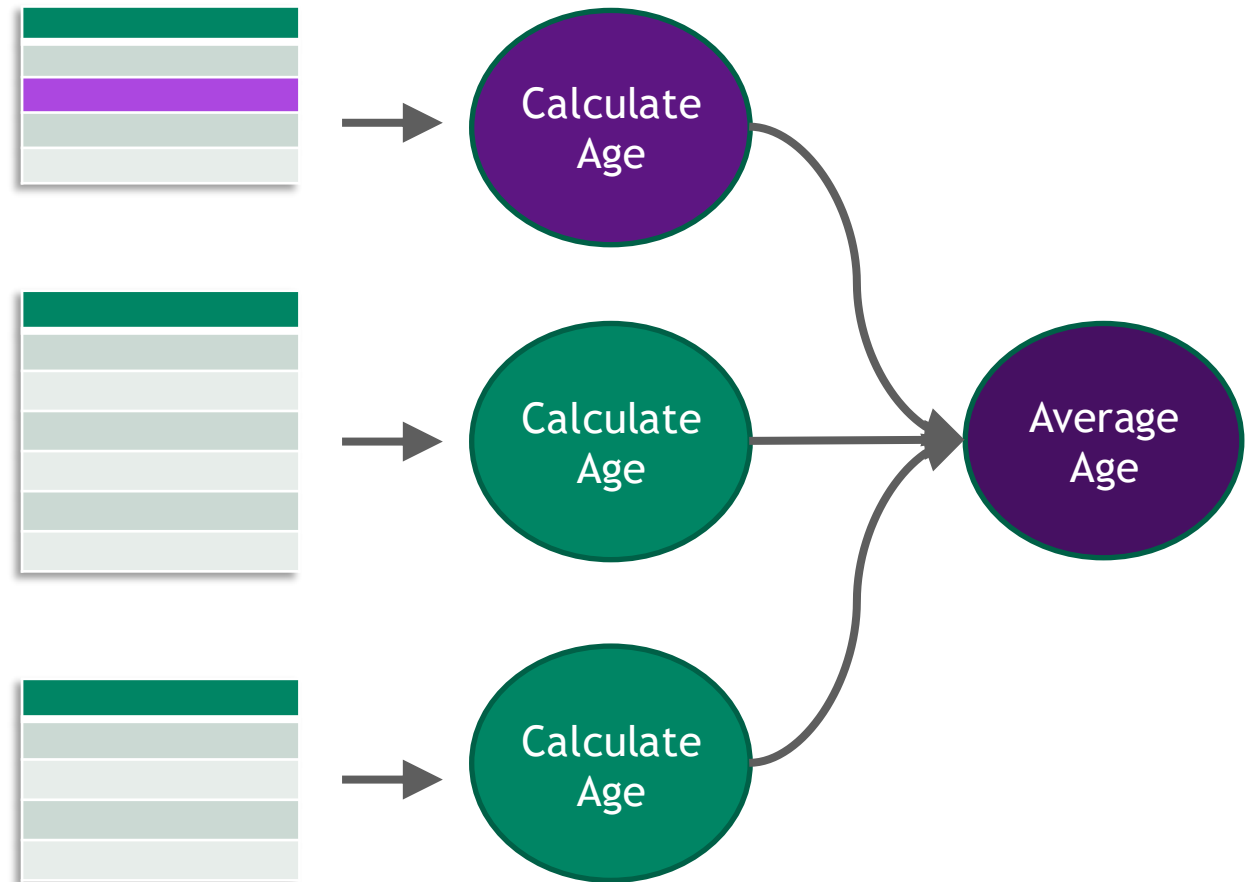
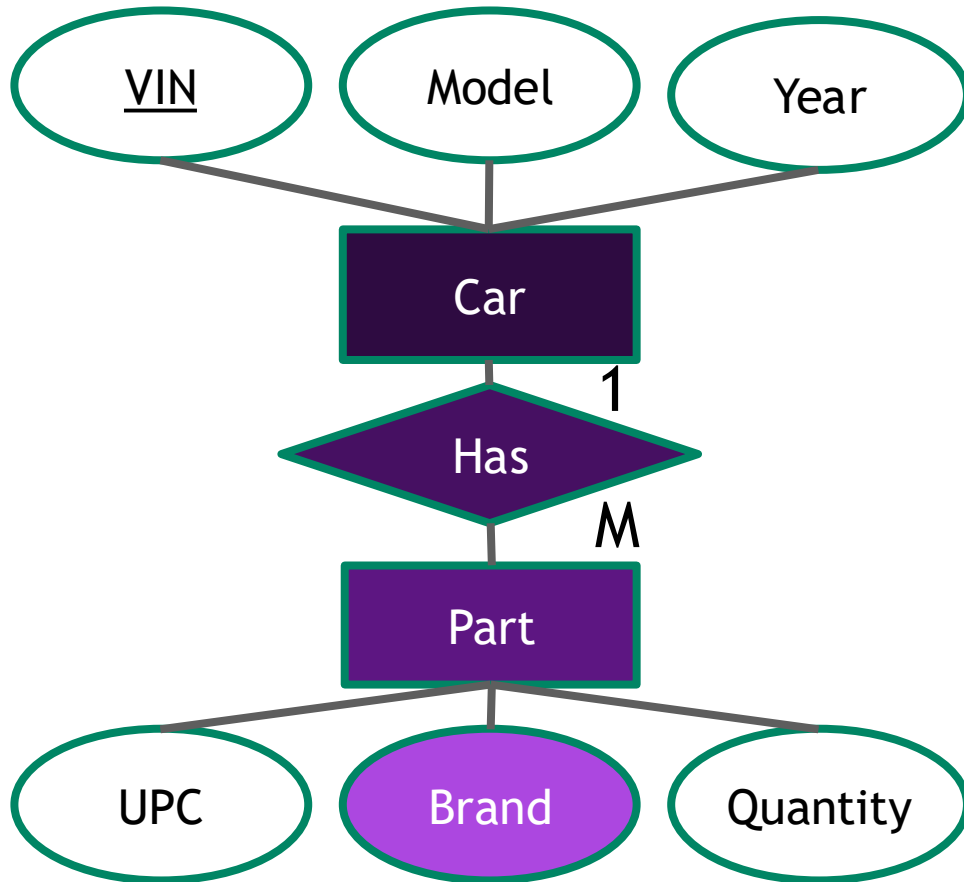
# DATABASE SYSTEM DESIGN

## Directed Acyclic Graphs



# DATABASE SYSTEM DESIGN

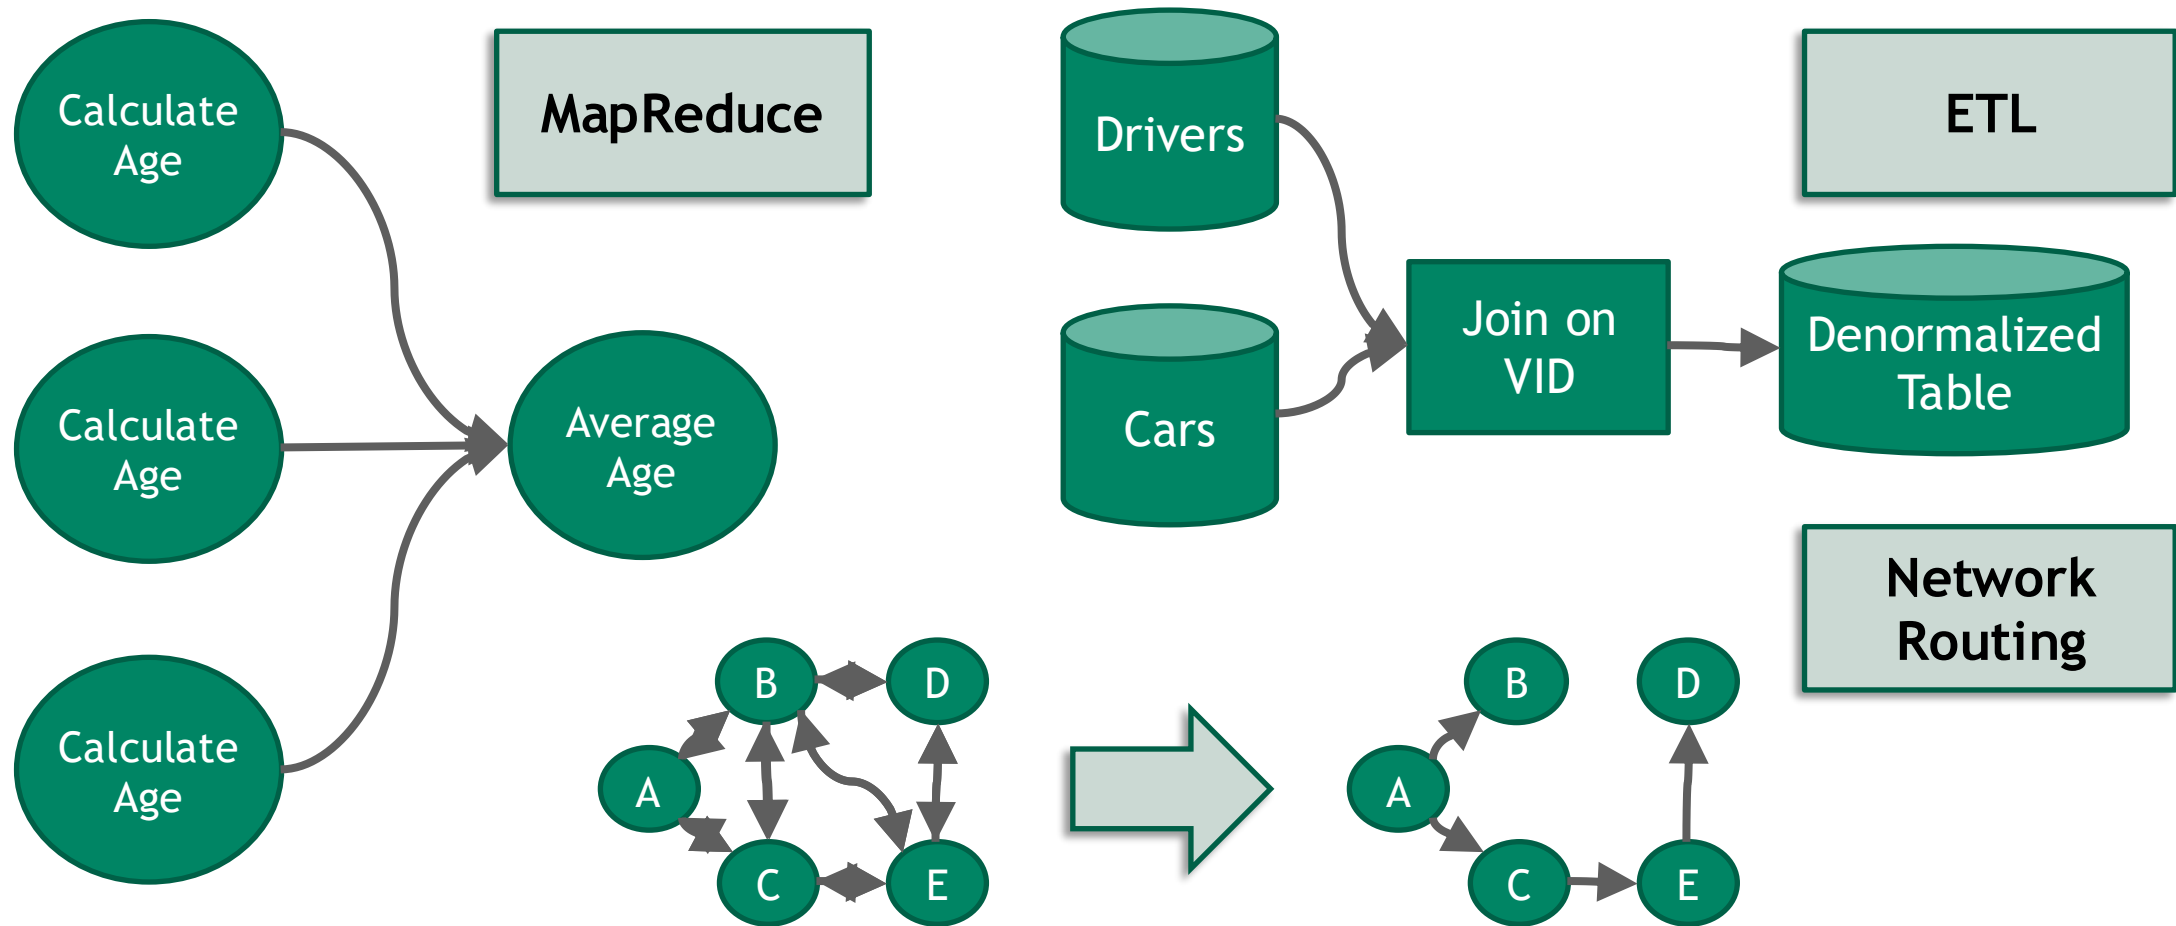
Data Quality





# DATABASE SYSTEM DESIGN

## Directed Acyclic Graphs

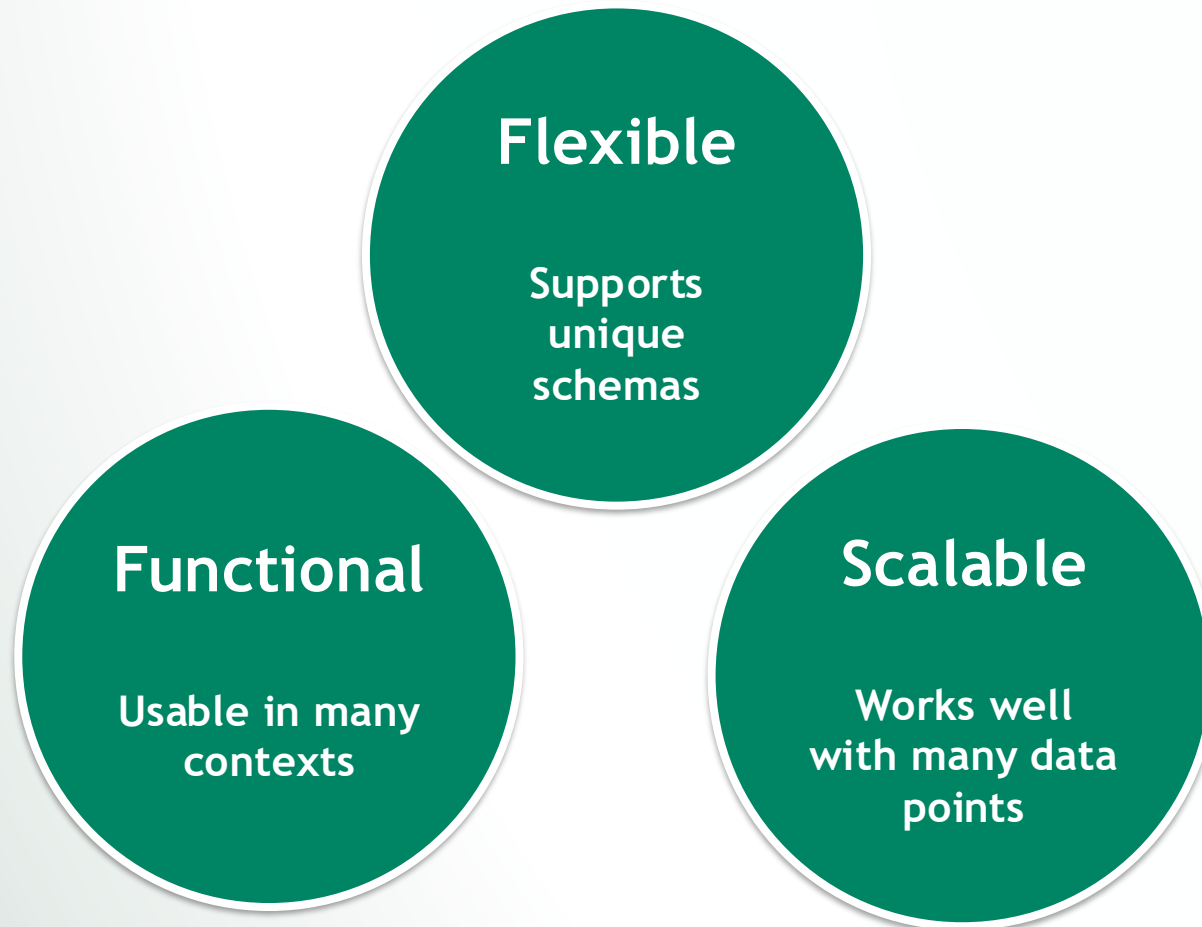




FILE FORMATS

# DATA FORMATS

Pick 1 - 2



These definitions vary based on context:

- Scalable to read or to write?
- Scalable with speed or cost?
- Flexible in the data store or flexible in the application?
- Functional for the server or functional for the client?

# PICKING THE BEST FORMAT

## CRUD

- Create
  - Add a record
- Read
  - Get record
- Update
  - Change a record
- Delete
  - Remove a record



# ROW VS COLUMNAR STORAGE

- Row -

Efficient for  
Adding a new record

- 
- Formats
    - CSV (Comma-Separated Values)
    - TSV (Tab-Separated Values)
    - Apache AVRO
  - Engines
    - MySQL
    - PostgreSQL

| Columnar |

Efficient for  
Data Aggregation

- 
- Formats
    - Apache Parquet
  - Engines
    - BigQuery
    - Snowflake
    - Redshift



# WRITING

## Adding a New Entry

### Row Formatted Data

Grace | Hopper | 1906

Blaise | Pascal | 1623

Katherine | Johnson | 1918

Alan | Turing | 1912

Can concatenate to end, or  
inserted by row number

First	Last	Born
Grace	Hopper	1906
Blaise	Pascal	1623
Katherine	Johnson	1918

+

Alan	Turing	1912
------	--------	------

### Column Formatted Data

Grace | Blaise | Katherine

Hopper | Pascal | Johnson

1906 | 1623 | 1918

Alan

Turing

1912

Broken up and inserted at  
the end of each block

# ANALYSIS

## A.K.A Feature Engineering

### Row Formatted Data

Grace | Hopper |

Blaise | Pascal |

Katherine | Johnson |

GH

BP

KJ

Broken up and inserted at  
the end of each block

First	Last	In.
Grace	Hopper	GH
Blaise	Pascal	BP
Katherine	Johnson	KJ
Alan	Turing	AT

>

### Column Formatted Data

Grace | Blaise | Katherine |  
Hopper | Pascal | Johnson |

GH | BP | KJ

Can be concatenated to end or  
inserted by column number

# BINARY

Ex: Multimedia File



Pro

- Compact
  - Faster to send and process
- Flexible
  - Many datatypes can easily be converted to binary
  - Great for images



Con

- Hard to visualize without decoding software
  - Difficult to debug data integrity


# ASCII

Ex: CSV



Pro

- Simple structure
  - No file metadata
- File is human readable
- Average scalability
  - Easy to join and split multiple CSV files
  - Easy to append a new entry



Con

- Simple structure
  - No file metadata
- Average scalability
  - Data is not compressed as much as other file types

# PARQUET

Ex: Hadoop



Pro

- Good compression if many repeated values
- Efficient to read a subset of columns
- Support for complex datatypes like arrays



Con

- Immutable
  - Query results are typically saved in a new file
- Querying for all the attributes of an entity is an expensive operation
- Files are not human readable without a tool



# DATA FORMATS COMPARISON

## Summary

Properties	CSV	JSON	Parquet	Avro
Columnar			✓	
Compressible	✓	✓	✓	✓
Splittable	✓	✓	✓	✓
Human readable	✓	✓		
Complex data structure		✓		
Schema evolution/validation		✓	✓	✓
Binary			✓	✓



# DATA FRAMEWORKS

# VERTICAL VS HORIZONTAL SCALING

↑ Vertical ↑

Scales to higher quality hardware

---

- SQL
- CuPy
- NumPy
- cuDF
- pandas

← Horizontal →

Scales to more partitions / machines

---

- Dask
- NoSQL
- Spark
- Hadoop

# SQL

## Structured Query Language

### Query

```
SELECT first, last  
FROM awesome.people  
WHERE born > 1900
```

### Table

(awesome.people)

first	last	born
Grace	Hopper	1906
Blaise	Pascal	1623
Katherine	Johnson	1918
Alan	Turing	1912

### Result

first	last
Grace	Hopper
Katherine	Johnson
Alan	Turing

# DATAFRAMES

Pandas (CPU) and cuDF (GPU)

Query



```
df = df[df["born"] > 1900]
```

```
df = df["first", "last"]
```

Table

df

first	last	born
Grace	Hopper	1906
Blaise	Pascal	1623
Katherine	Johnson	1918
Alan	Turing	1912



first	last
Grace	Hopper
Katherine	Johnson
Alan	Turing

Result



# MATRICES AND NUMBER ARRAYS

NumPy (CPU) and CuPy (GPU)

Query



```
a = a.sum(axis = 0)
```

Array

a

8	6	7
5	3	0
9	8	6
7	5	3

Result

29

22

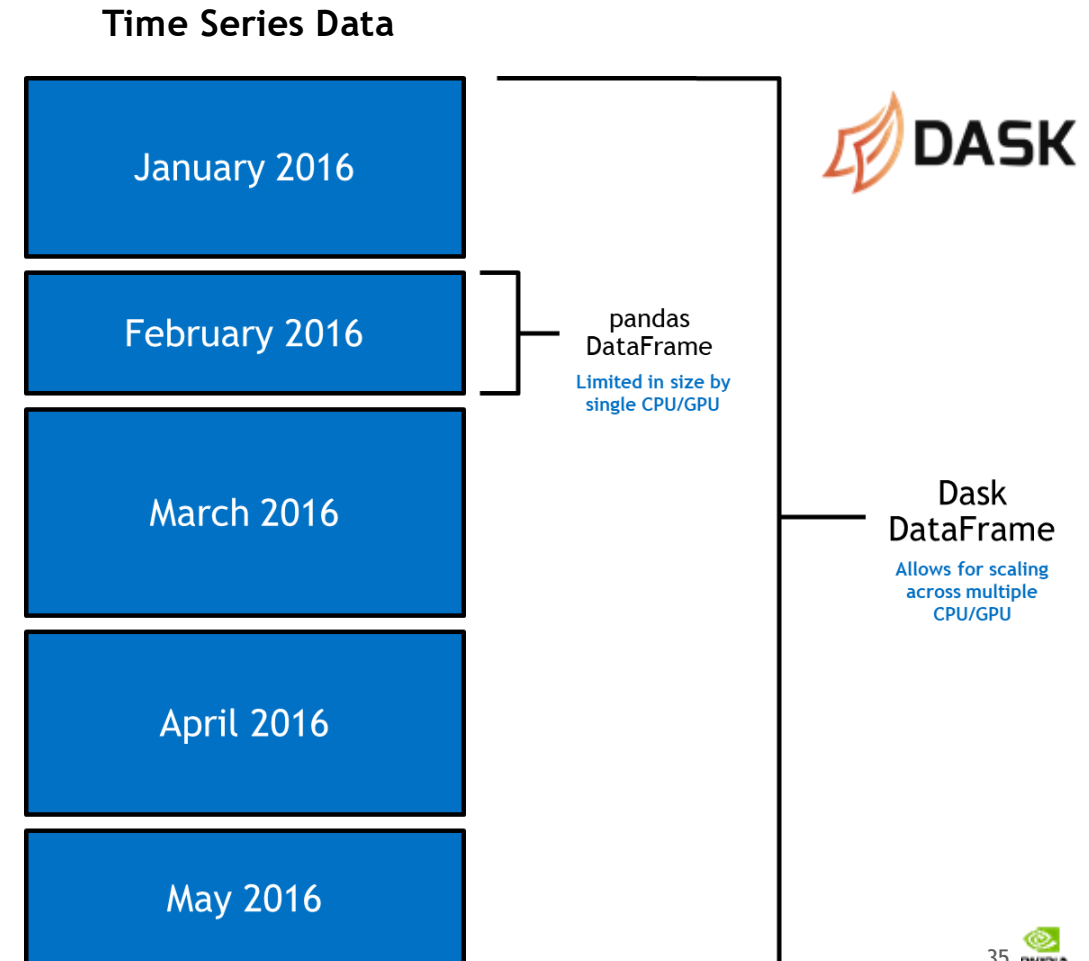
16



# DASK SCALES PYTHON ANALYTICS

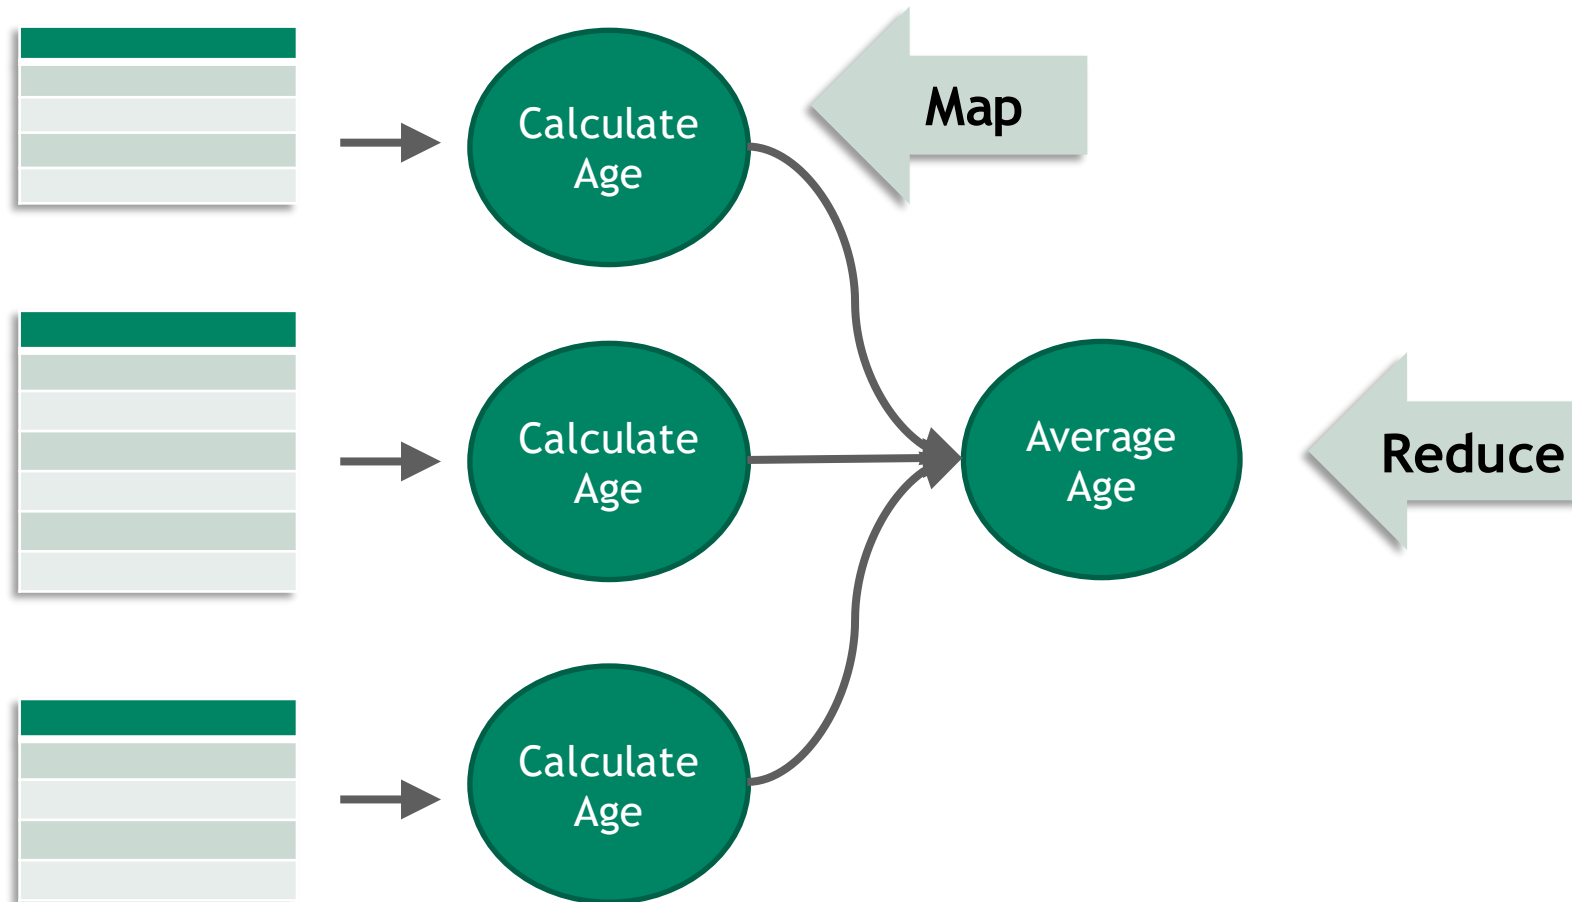
## SCALE FROM A LAPTOP TO LARGE-SCALE CLUSTERS WITH EASE

- **Dask** enables data scientists to scale out analytics workloads in native Python. With an optimized scheduler, Dask makes it easy to schedule and execute tasks on distributed computation.
- **Dask** follows the standards set by the PyData ecosystem to provide a familiar, comfortable user experience at scale. When paired with NVTABULAR/RAPIDS, data scientists can leverage the processing power of NVIDIA accelerated compute and distribute across clusters to improve cycle time-reducing time to insights drastically.



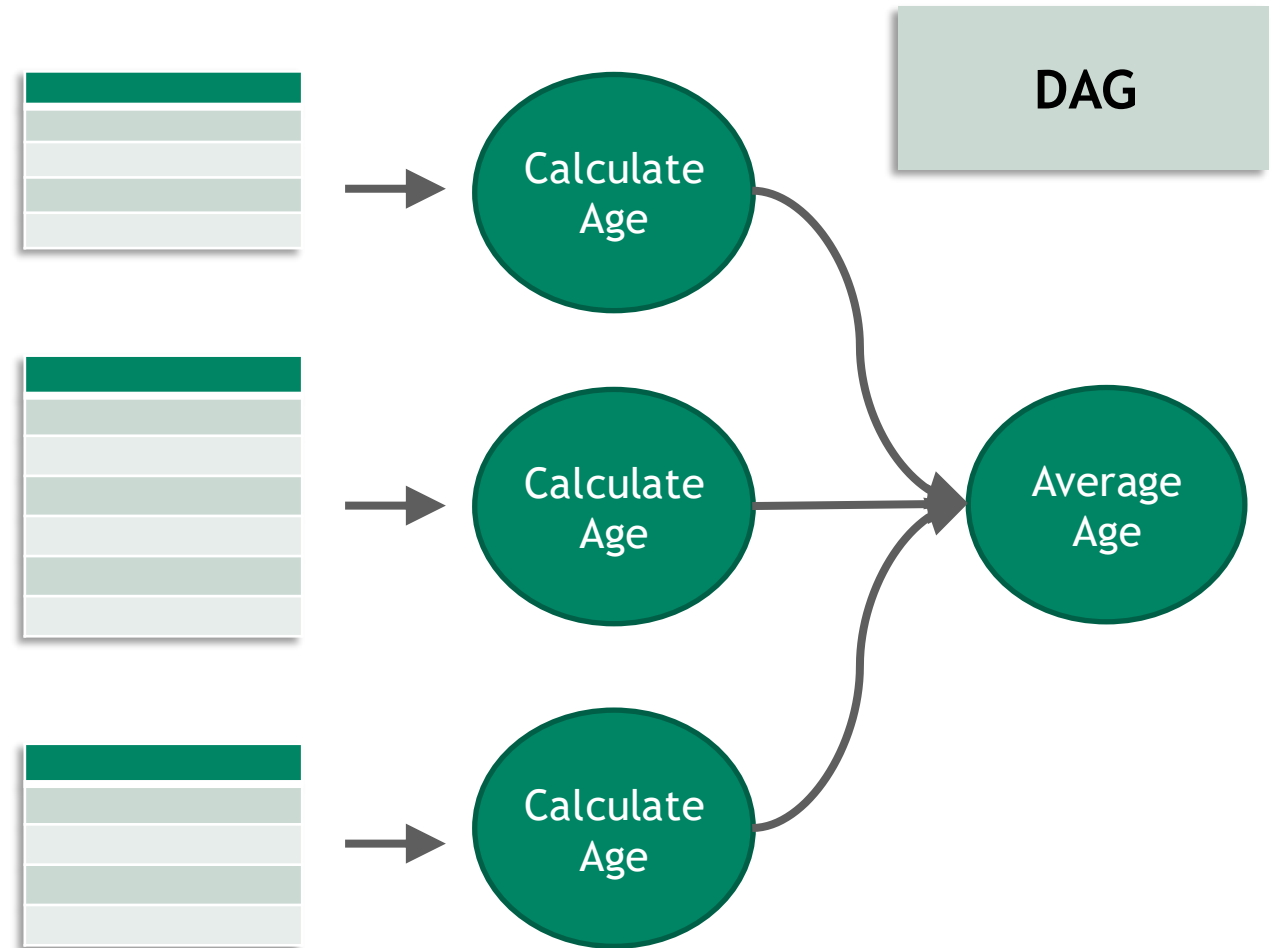
# MAPREDUCE

Map to each thread, Reduce all threads to one



# LAZY EXECUTION

## Building a Factory



# RELATIONAL DATABASES

Ex: SQL



Pro

- Well known
- Concise Language
- Relatively fast querying
  - Foreign keys
- Blazing SQL



Con

- Inflexible data structure
  - Some objects do not convert well to table format
- Typically, single server
  - More expensive hardware needed to scale

# DATAFRAME

Ex: cuDF, Pandas, R



Pro

Python and R APIs

- cuDF, Pandas
- Compared to SQL, more flexible operations
- Easier to make user-defined functions and integrate third party libraries



Con

- Single server, not meant for large-scale data manipulation
  - Consider Spark instead
- Compared to SQL, not as scalable

# DASK

Ex: Dask DataFrame, Dask-cuDF



Pro

- Large computation can receive a significant speed increase
- Can read large data sources due to partitioning



Con

- Large overhead to set up not worth it for small files or limited computation
- Lazy execution can make it tricky to debug





LAB


# WEATHER SYSTEMS



Credit: Ralph F. Kresge, Submitted to NOAA


# INVESTIGATING WATER LEVEL

NOAA/NOS/CO-OPS  
Observed Water Levels at 8735523, East Fowl River Bridge AL  
From 2021/03/10 00:00 GMT to 2021/03/11 23:59 GMT





LET'S GO!



DEEP  
LEARNING  
INSTITUTE